

EE 8310

Assignment - 3

AI Agent for SPICE Automation

Nandini Kumawat
kumaw010@umn.edu

The Challenge

Time-consuming netlist creation and parameter sweeps

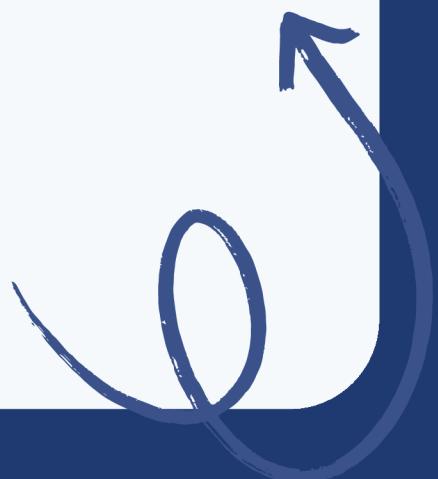
Manual SPICE simulations are **labor-intensive** processes, requiring extensive time spent on netlist creation and parameter sweeps, hindering efficient circuit design and delaying project timelines significantly.

- Natural-language → parameter parser → HSPICE netlist generator
- Batch mode: auto-measurements, CSV logs
- Interactive mode: waveform viewing in ngspice GUI
- Targets: Inverter & NAND2 (VDD / Temp / Cload sweeps)

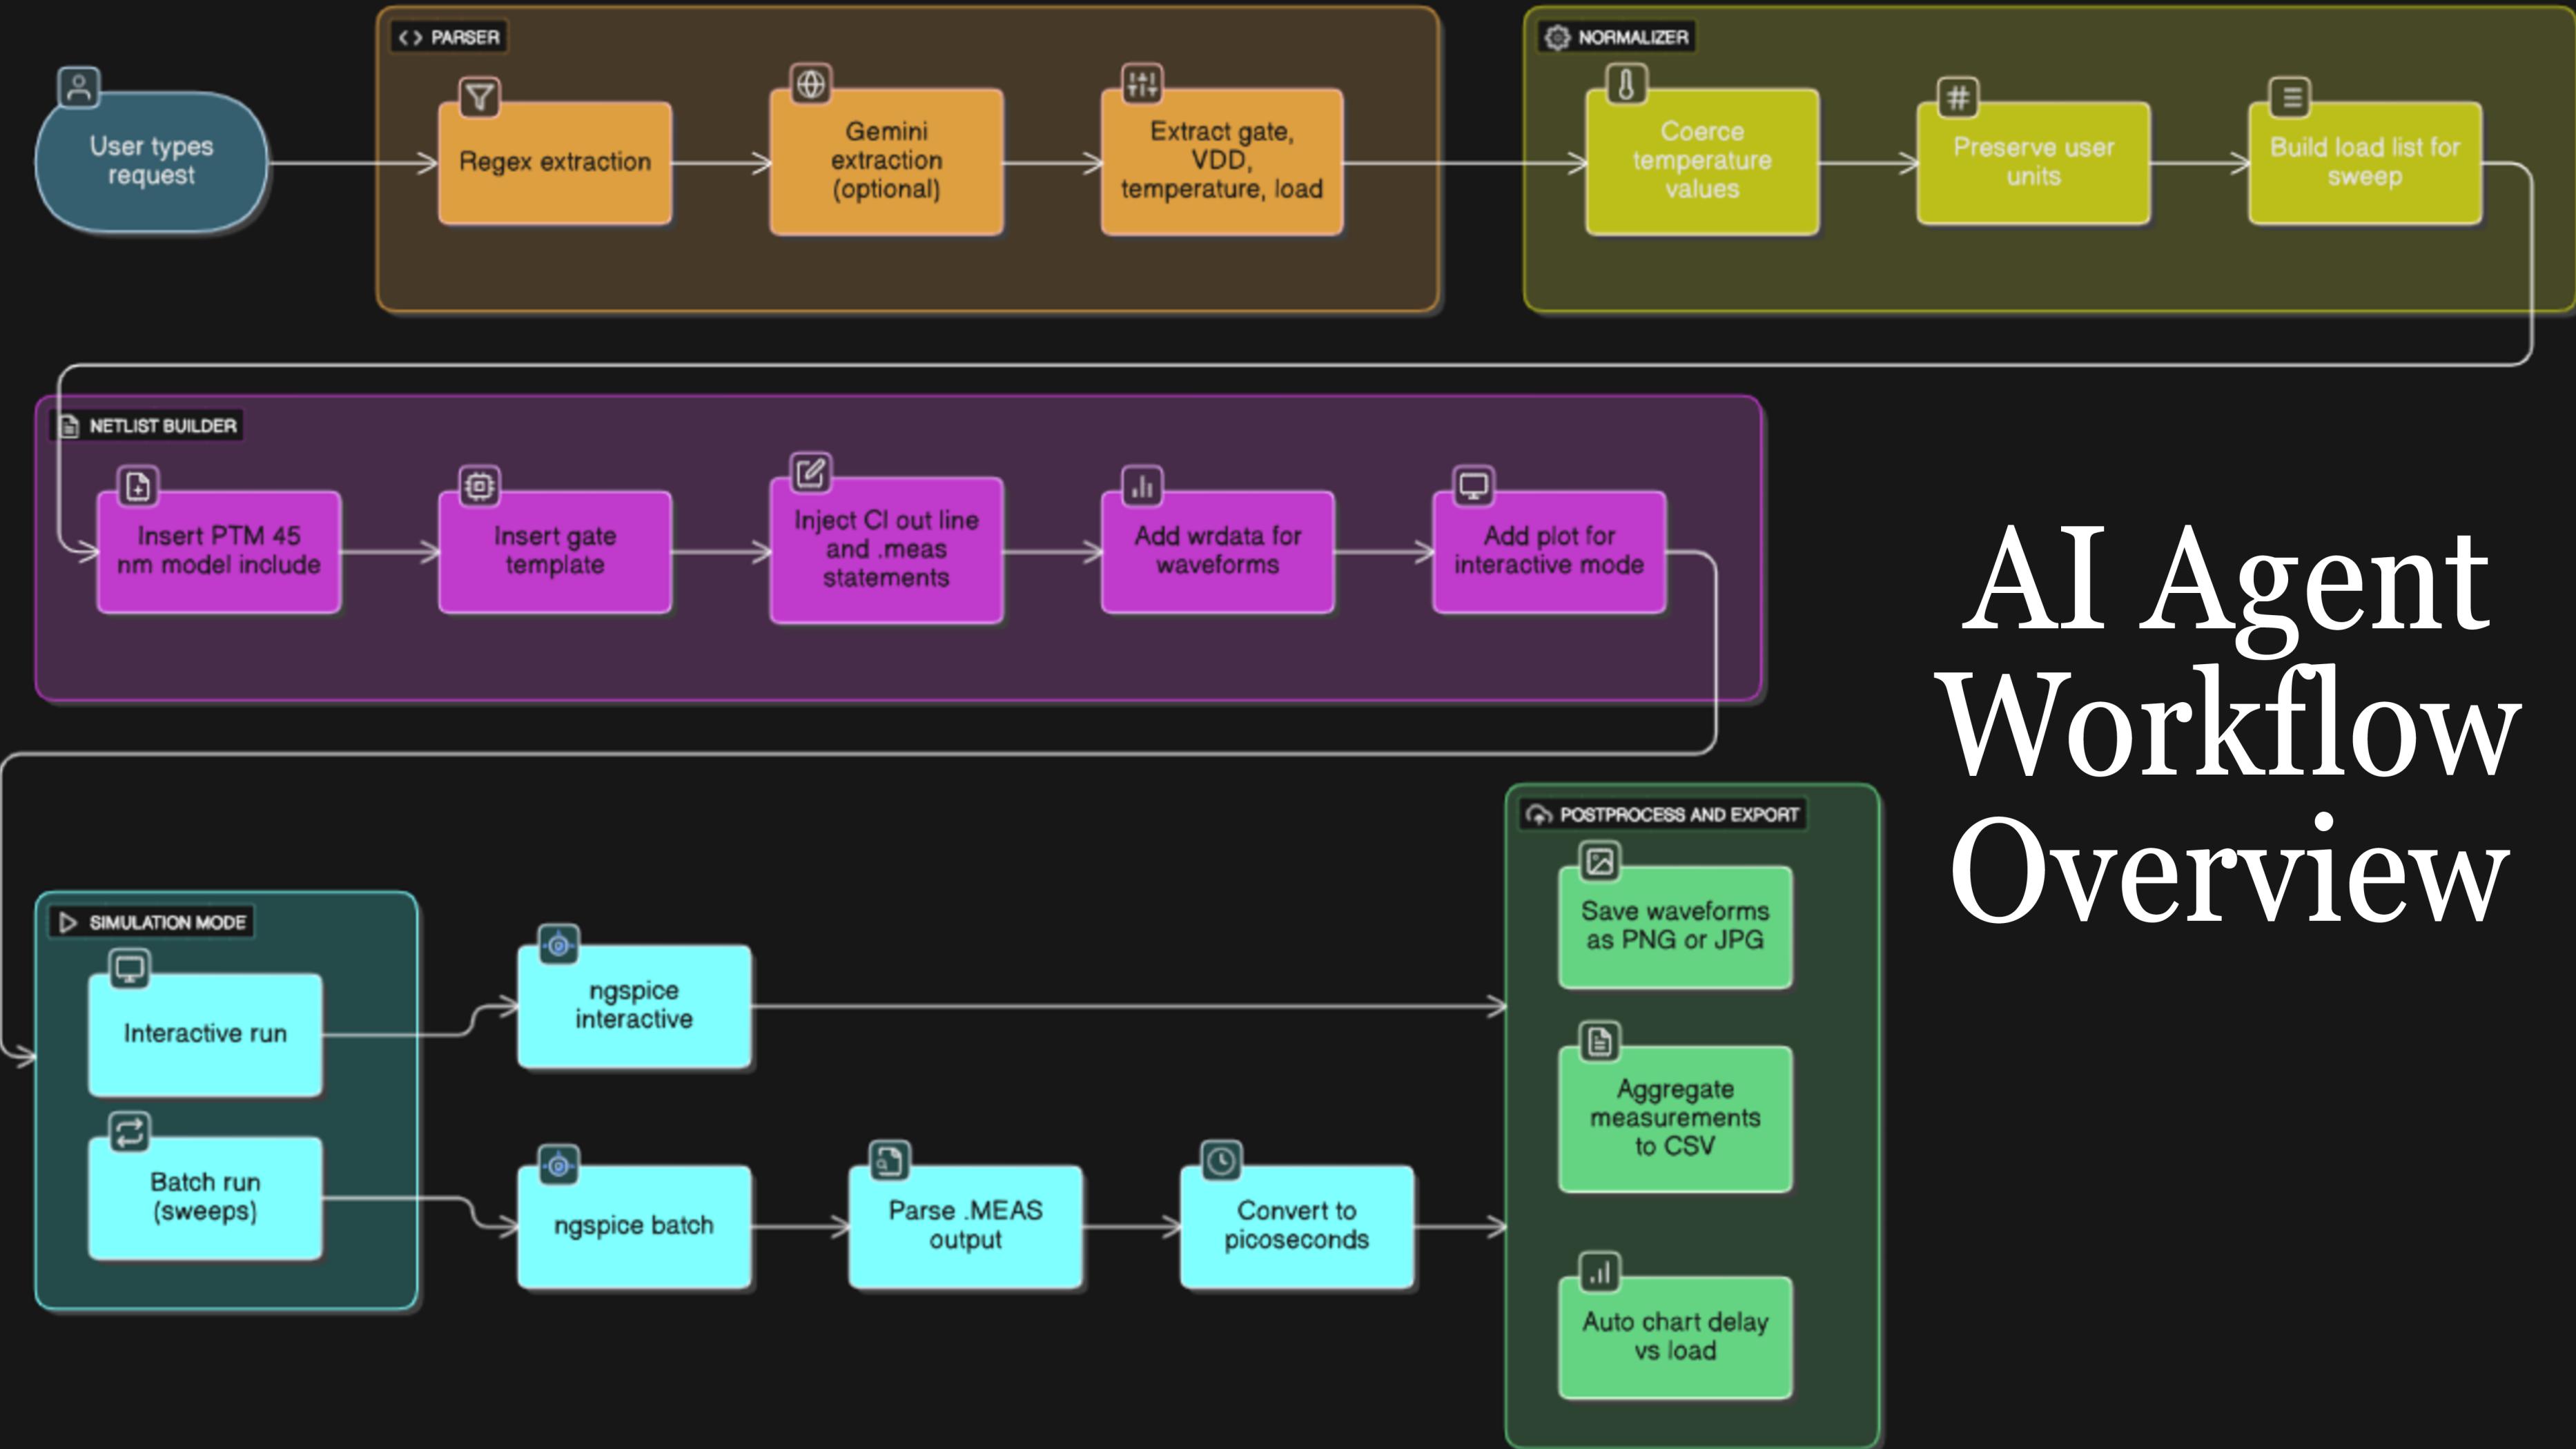
Solution

My agent can:

- i. Accept NL requests (e.g., “nand2, vdd 0.72, load 5–50 fF step 5 fF, temp 25C”).
- ii. Parse VDD / Temp / Cload / Gate type via regex rules (optionally Gemini).
- iii. Write ngspice netlists dynamically with .tran, .measure, .print/.probe, .option post.
- iv. Run ngspice batch to extract .mt0 measurement tables and .tr0 waveforms; aggregate to CSV.
- v. View waveforms (interactive): open .tr0 in ngpsice GUI for waveform plots



AI Agent Workflow Overview



User Prompts

- “inverter vdd 0.8, temp 25 C, load 5–50 fF step 5 fF”
- → gate=inverter, VDD=0.8, Temps=[25], Loads=[5...50 fF]
- “nand2 at vdd 0.72, 0.8, 0.88 load 10 fF temp 27 C”
- → gate=nand2, VDD sweep=[0.72,0.8,0.88], Temp=27, Load=10 fF
- “inverter temps -40, 25, 110C vdd 0.8 load 15 fF”
- → Temp sweep, fixed VDD & Cload

CHEATSHEET

```
prompt    := gate_clause (vdd_clause | temp_clause | load_clause)*
gate_clause := ( "gate" | cellname ) cellname
vdd_clause := ("vdd" | "at vdd") values
temp_clause := ("temp" | "temps" | "temperature") values
load_clause := ("load" | "cap" | "Cload" | "capacitance") values

values    := list | range | scalar
list     := scalar ("," scalar)+
range    := scalar ("-"|"--") scalar "step" scalar
scalar   := number [unit]
unit     := "V" | "mV" | "C" | "fF" | "pF" | "nF"
```

```
nandini@DESKTOP-781T80K:/mnt/e/University of Minnesota/Fall '25/EE8310 Advanced Topics in VLSI/Project 3$ export GEMINI_API_KEY="AIzaSyBJDE7awTvPuzQ4l8L88ksa1ZgHUlkbMa0"
nandini@DESKTOP-781T80K:/mnt/e/University of Minnesota/Fall '25/EE8310 Advanced Topics in VLSI/Project 3$ python3 ai_spice_agent.py python3 ai_spice_agent.py --mode batch
usage: ai_spice_agent.py [-h] [--mode {batch,interactive}] [--open-images] [--jpg]
ai_spice_agent.py: error: unrecognized arguments: python3 ai_spice_agent.py
nandini@DESKTOP-781T80K:/mnt/e/University of Minnesota/Fall '25/EE8310 Advanced Topics in VLSI/Project 3$ python3 ai_spice_agent.py --mode batch
Enter your simulation request: inverter vdd 0.8 V, temp 25 C, load 5-50 fF step 5 fF
```

Propagation Delay Trends vs VDD, Temperature, and Load

Inverter @ VDD = 0.8 V (fixed input edge)

Temperature sweep (single-point load)

- -40°C : $\text{tplh} \approx 12.43 \text{ ps}$, $\text{tphl} \approx 13.66 \text{ ps}$
- 25°C : $\text{tplh} \approx 16.76 \text{ ps}$, $\text{tphl} \approx 18.46 \text{ ps}$
- 110°C : $\text{tplh} \approx 24.08 \text{ ps}$, $\text{tphl} \approx 26.57 \text{ ps}$
- Trend: delay roughly doubles from $-40^{\circ}\text{C} \rightarrow 110^{\circ}\text{C}$.

Load sweep (25°C , VDD 0.8 V, $C = 5 \rightarrow 50 \text{ fF}$, step 5 fF)

- tplh : $29.03 \rightarrow 132.73 \text{ ps} \Rightarrow \sim 2.30 \text{ ps/fF}$ slope
- tphl : $32.57 \rightarrow 149.50 \text{ ps} \Rightarrow \sim 2.60 \text{ ps/fF}$ slope
- Trend: near-linear with C_{load} ; falling is consistently slower.

Propagation Delay Trends vs VDD, Temperature, and Load

NAND2, Cload = 5 fF, T = 27 °C (VDD sweep)

- VDD = 0.72 V: t_{phl} ≈ 120.73 ps, t_{plh} ≈ 40.37 ps
- VDD = 0.80 V: t_{phl} ≈ 68.33 ps, t_{plh} ≈ 29.42 ps
- VDD = 0.88 V: t_{phl} ≈ 47.03 ps, t_{plh} ≈ 21.89 ps
- Trend: both edges speed up with VDD; t_{phl} shows stronger VDD sensitivity (series NMOS discharge path).

NAND2, VDD = 0.8 V, Cload = 10 fF (Temp sweep)

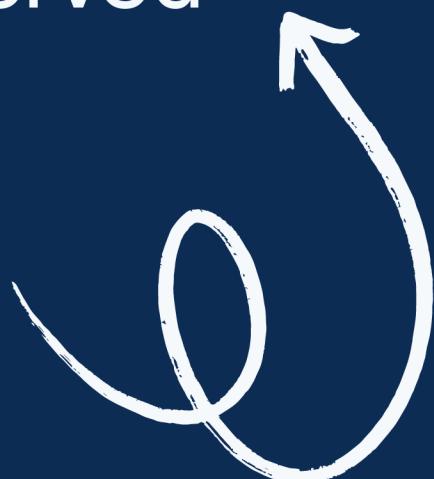
- -40 °C: t_{phl} ≈ 48.61 ps, t_{plh} ≈ 24.10 ps
- 25 °C: t_{phl} ≈ 94.88 ps, t_{plh} ≈ 39.24 ps
- 110 °C: t_{phl} ≈ 169.17 ps, t_{plh} ≈ 29.54 ps
- Trend: t_{phl} rises strongly with Temp (NMOS mobility loss dominates). t_{plh} is non-monotonic here (likely PMOS vs NMOS temp-V_{th} interplay and exact slew/measurement points); still, rising edges remain faster than falling for this setup.

Propagation Delay Trends vs VDD, Temperature, and Load

- $V_{DD} \uparrow \rightarrow \text{delay} \downarrow$ (strong effect)
- $\text{Temp} \uparrow \rightarrow \text{delay} \uparrow$ (mobility \downarrow ; V_{th} shifts)
- $C_{load} \uparrow \rightarrow \text{delay} \uparrow \sim \text{linearly}$ (RC charging/discharging)
- Falling vs rising: typically $t_{phl} > t_{plh}$ (drive asymmetry); NAND2 shows stronger asymmetry due to series NMOS stack.

Analysis – How the AI agent helps

- Manual effort reduced vs hand-writing netlists
- Natural-language to deck: “inverter vdd 0.8 temp -40,25,110 load 15 fF” → complete ngspice deck with .tran, .measure, .probe (no boilerplate typing).
- Param sweeps in one go: temps/VDD/Cload lists & ranges expand automatically into runs; results collated into CSV and delay-vs-load/Temp/VDD plots.
- Consistent measurement windows: trigger/target rules (50% VDD, TD guards) are generated uniformly → fewer “missed target” mistakes.
- One-click waveforms: batch writes .tr0 and opens GUI for viewing; no manual file hunting.
- Reproducibility: the same prompt regenerates the same deck(s) with preserved units (e.g., 5fF, not 5e-15), easing reviews and graded checks.
-



Limitations / risks with AI automation

- Hallucinated parameters: LLMs can invent defaults (e.g., pulse widths, device sizes). We mitigated by rules-first parsing and only optional LLM assist.
- Syntax & dialect drift: small differences across PDKs/BSIM options and tool versions (HSPICE vs ngspice) can break runs; templates must be curated.
- Parser brittleness: odd phrasing or unit forms (“5 fF step 5 fF to 50 fF”) needed careful regex handling; we added range parsing and unit canonicalization.
- Debugging overhead: when a run fails, you still read simulator logs (.lis, .mt0), fix models/paths, or adjust .measure; AI doesn’t replace EDA debugging skills.
- Result trust: automation is fast, but engineers must sanity-check measurements and waveforms (e.g., ensure the first valid edge is measured, stacking cases in NAND2).

Scaling to larger designs

- Structured, hierarchical builds. Pull subcircuits from libraries, auto-wire the testbench, and apply a standard .measure bundle across every block—clean, repeatable, and scalable.
- Corners at scale. Generate PVT grids and Monte Carlo runs programmatically, log to CSV/Parquet, and feed dashboards for delay, leakage, and yield—so coverage is comprehensive, not ad hoc.
- Policy + lint before compute. Run pre-flight checks for floating nodes, missing supplies, and min/typ/max model selection to catch issues before expensive simulations.
- Closed-loop exploration. Read trends (delay vs. C, VDD, T) and automatically propose the next most informative sweep—an active-learning–style DOE that accelerates convergence.



Pitfalls I actually hit (and fixes)

- Unit preservation: tool kept rewriting 5 fF as 5e-15.
- Fix: a formatter that echoes the user's units into the deck (prints 5fF exactly) and separate numeric conversion only for plotting/CSV.
- Scientific notation in .measure output (e.g., 2.9e-11):
- Fix: convert to picoseconds both inside the simulator (derived variables) and on the Python side for CSV/print.
- Range parsing crash (int() of dict in temp sweep):
- Fix: a coercion helper that turns strings, dicts, or lists into clean integer temperatures; protects against gemini/JSON quirks.
- GUI visibility: early runs didn't show waveforms.
- Fix: add .option post + .probe to guarantee .tr0; open those files explicitly.
- NAND2 measurements missing when both inputs toggled together:
- Fix: pulse only in1, hold in2 steady (VDD) so .measure always sees a valid targ; add TD margins.
- CSV wasn't showing up:
- Fix: central sweep_and_export path that always writes meas_sweep.csv and delay-vs-C plots after batch runs.
- “vin: no DC value” warning:
- Note: benign for PULSE sources in transient; we documented it to avoid confusion.

What I learned

01

- AI excels at the “glue.” It turns a plain request—“sweep 5–50 fF”—into consistent, linted decks with correct .measure blocks and auto-collected results. That replaces hours of copy-paste and human error.
- Guardrails are non-negotiable. Regex/rule parsing, unit preservation, fixed 50% trigger/target policies, and log parsing keep the flow deterministic and grading-safe. The LLM assists—it doesn’t dictate.
- Human circuit insight is still the core. Choosing TD windows, setting NAND biasing, and reading waveforms/.mt0 tables all require engineering judgment.
- Scale comes from structure. With the pipeline in place, adding corners, larger blocks, or new metrics is incremental—AI simply executes more experiments, more consistently, and much faster.

03